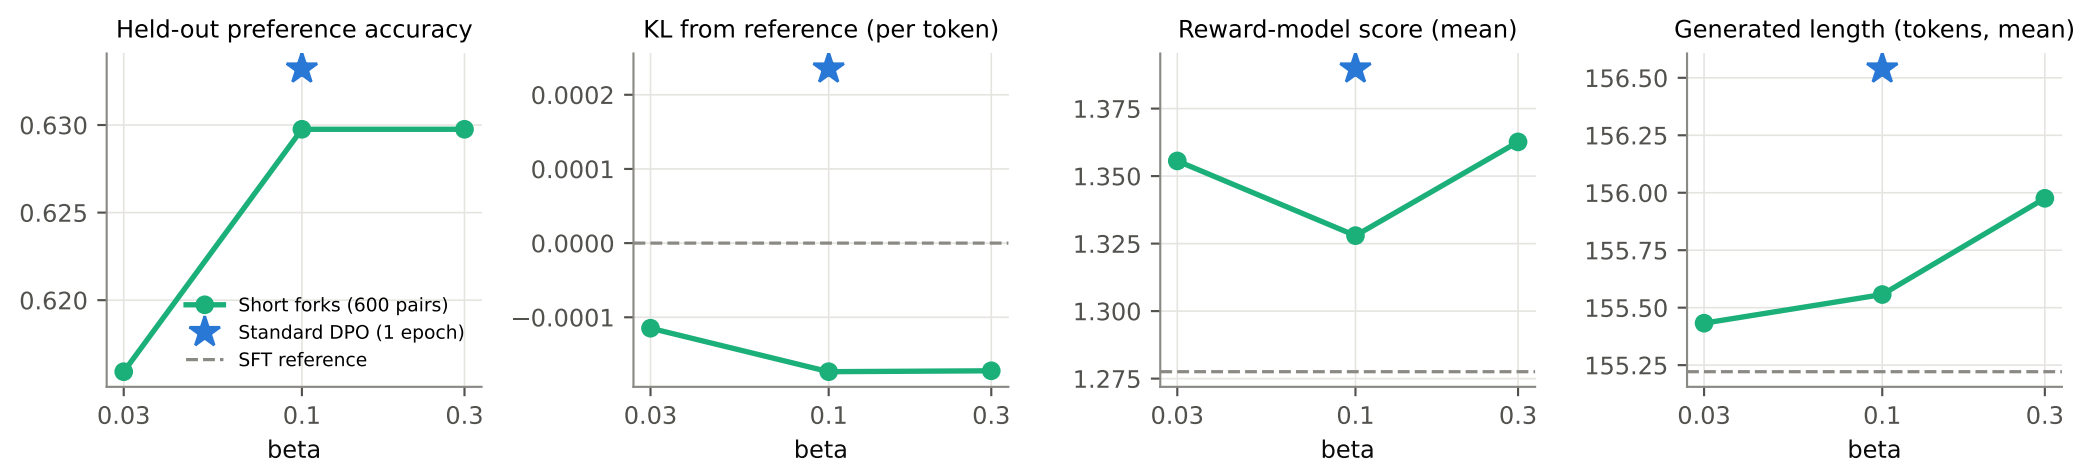
Held-out preference accuracy
0.630
0.625
0.620
Short forks (600 pairs)
Standard DPO (1 epoch)
SFT reference
0.03
0.1
0.3
beta
KL from reference (per token)
0.0002
0.0001
0.0000
−0.0001
0.03
0.1
0.3
beta
Reward-model score (mean)
1.375
1.350
1.325
1.300
1.275
0.03
0.1
0.3
beta
Generated length (tokens, mean)
156.50
156.25
156.00
155.75
155.50
155.25
0.03
0.1
0.3
beta