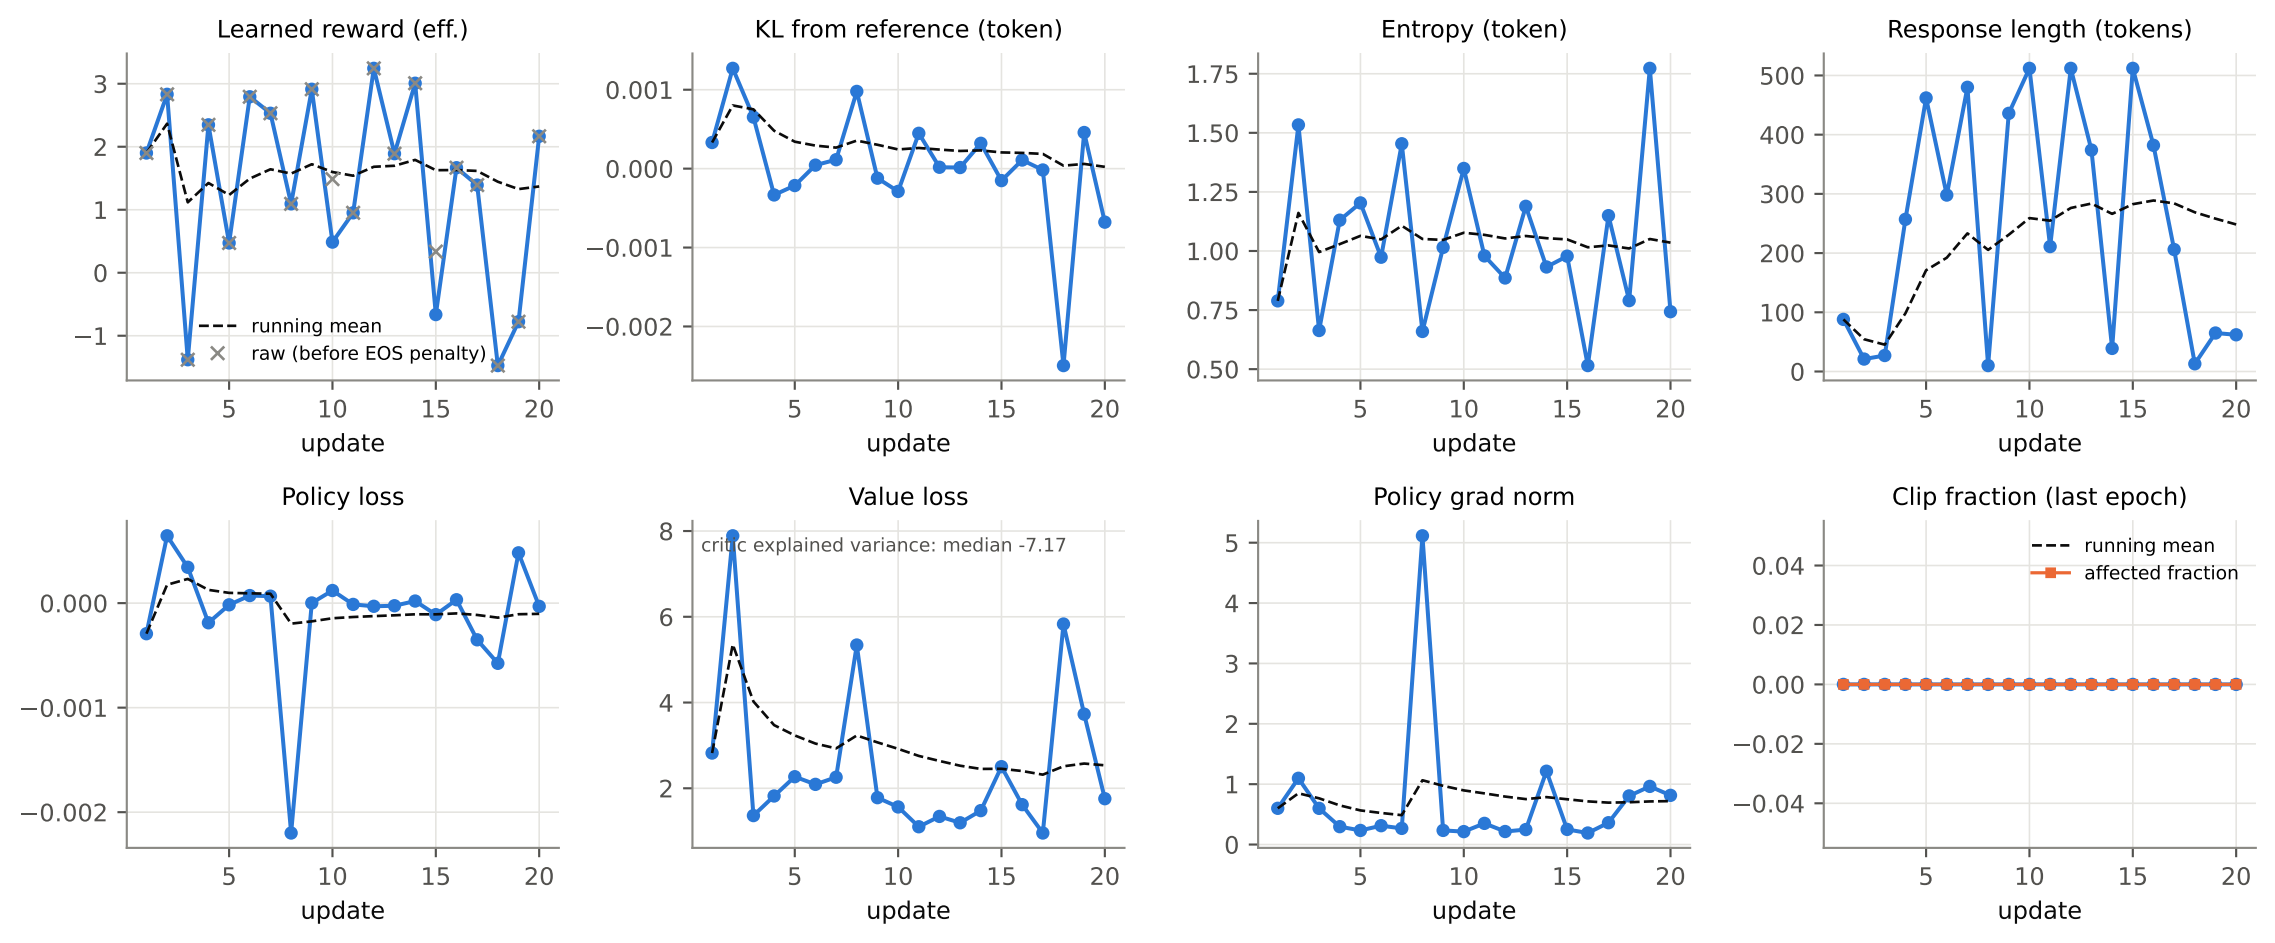
Learned reward (eff.)
running mean
raw (before EOS penalty)
update
KL from reference (token)
update
Entropy (token)
update
Response length (tokens)
update
Policy loss
update
Value loss
critic explained variance: median -7.17
update
Policy grad norm
update
Clip fraction (last epoch)
running mean
affected fraction
update