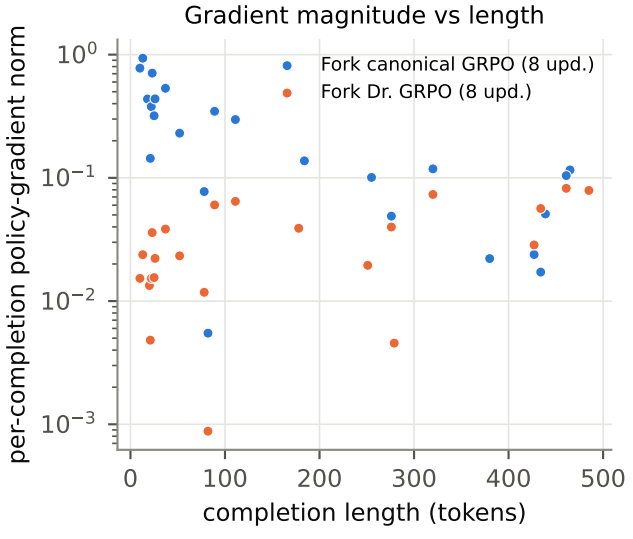
Gradient magnitude vs length
Fork canonical GRPO (8 upd.)
Fork Dr. GRPO (8 upd.)
per-completion policy-gradient norm
completion length (tokens)

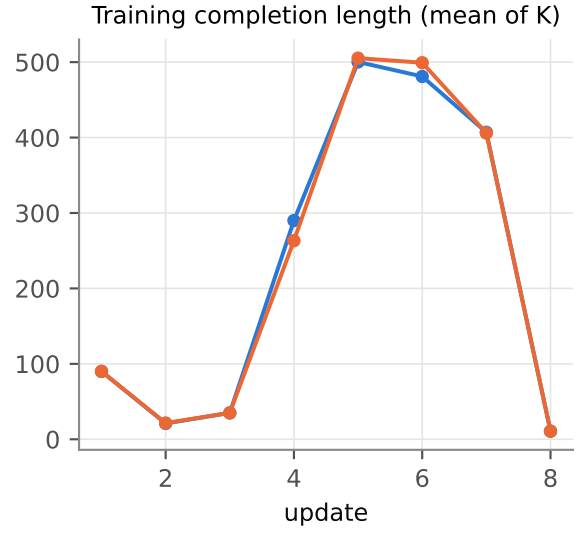
Training completion length (mean of K)
update

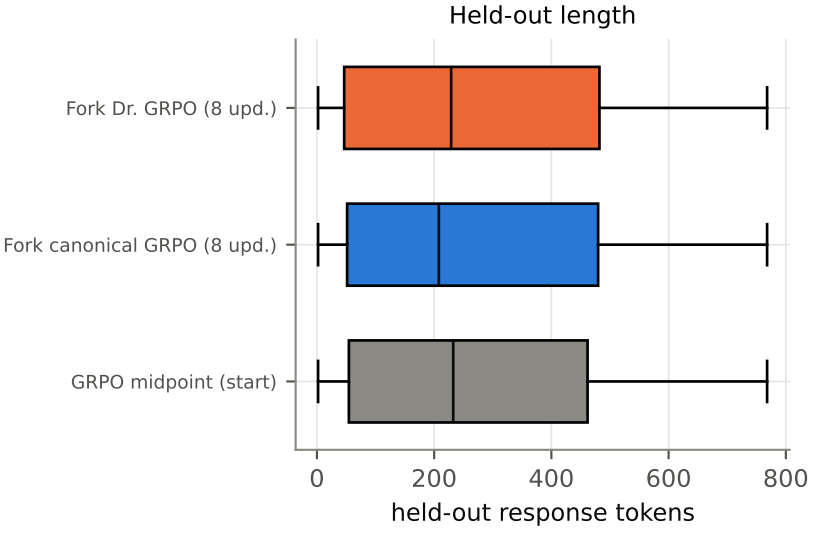
Held-out length
Fork Dr. GRPO (8 upd.)
Fork canonical GRPO (8 upd.)
GRPO midpoint (start)
held-out response tokens